

OBJECT DETECTION WITH YOLO

• • •

Michał Gilski and Boda Wen

YOLO

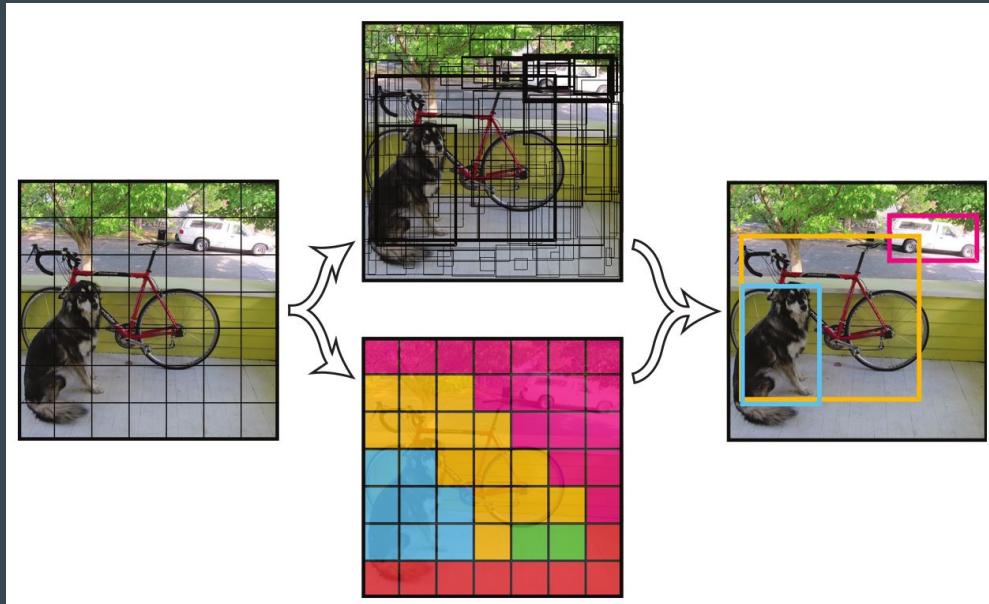
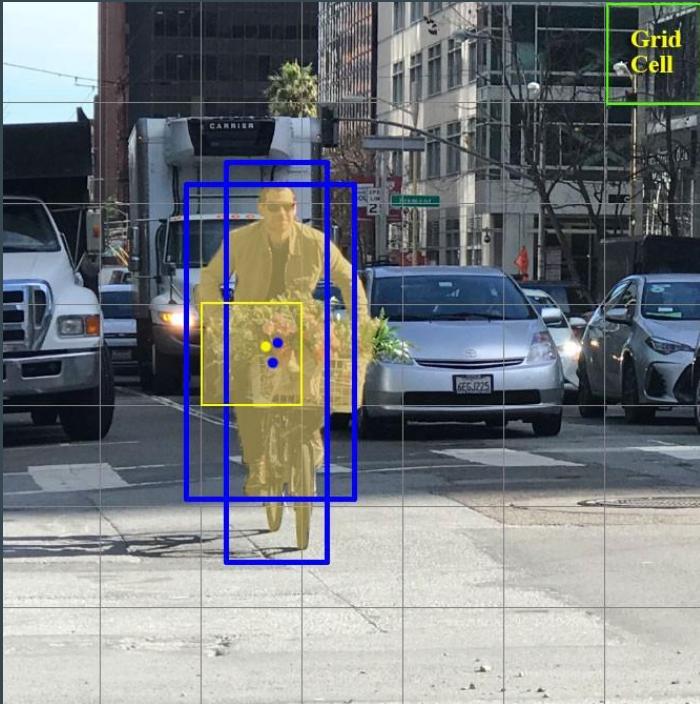
- YOLO (You Only Look Once)-> object detection
- DarkNet
- Earlier detection frameworks, looked at different parts of the image multiple times at different scales and repurposed image classification technique to detect objects. This approach is slow and inefficient.
- Region-based CNNs (R-CNNs):An object detector that has candidate bounding boxes that could contain objects. These regions were then passed into a CNN for classification.
- YOLO takes entirely different approach. It looks at the entire image only once and goes through the network once and detects objects. Hence the name. It is very fast. That's the reason it has got so popular.

Yolo history

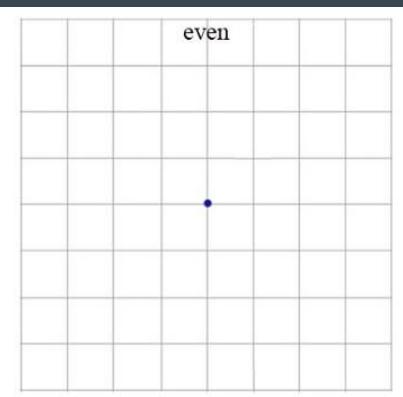
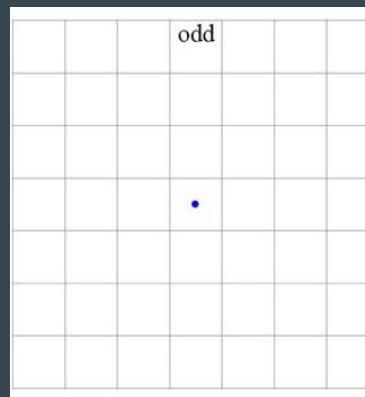
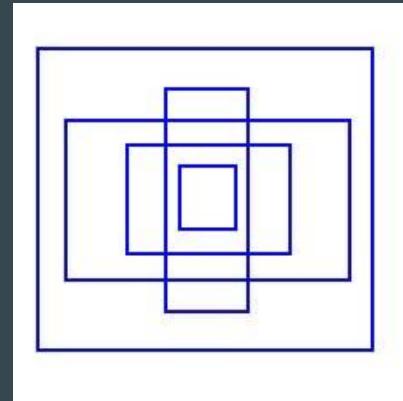
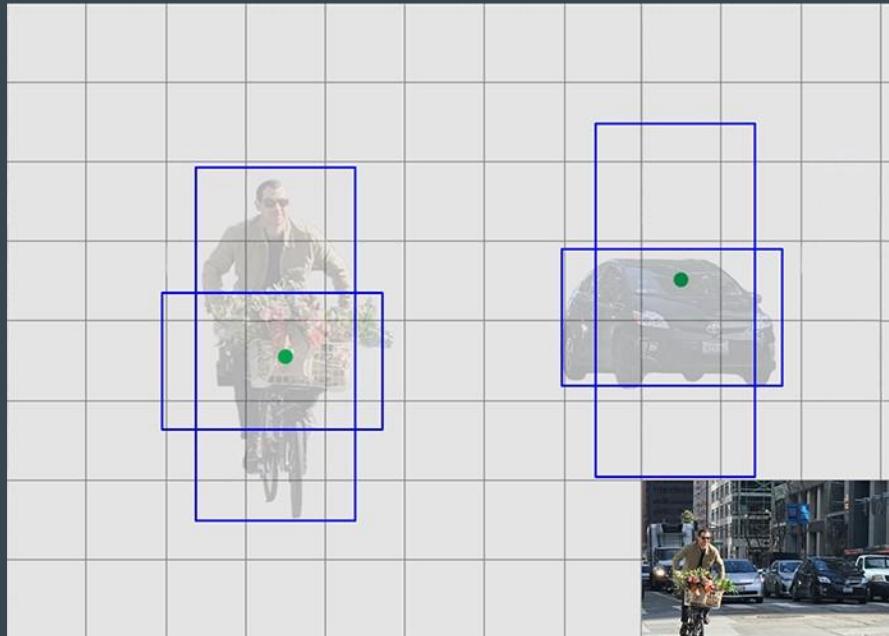
Created by Joseph Redmon and Ali Farhadi from University of Washington

- The first YOLO - faster than any other detection approach at 55 FPS
- Fast YOLO - more shallow network, but up to 155 FPS
- YOLOv2 - **Better, Faster, Stronger** (Batch normalization, introduction of Darknet and anchor boxes, higher resolution, Word trees, etc.)
- YOLO 9000 - detects over 9000 classes
- YOLOv3 - “**We mostly took good ideas from other people**”: three output layers, deeper network with shortcuts, feature extraction, etc..

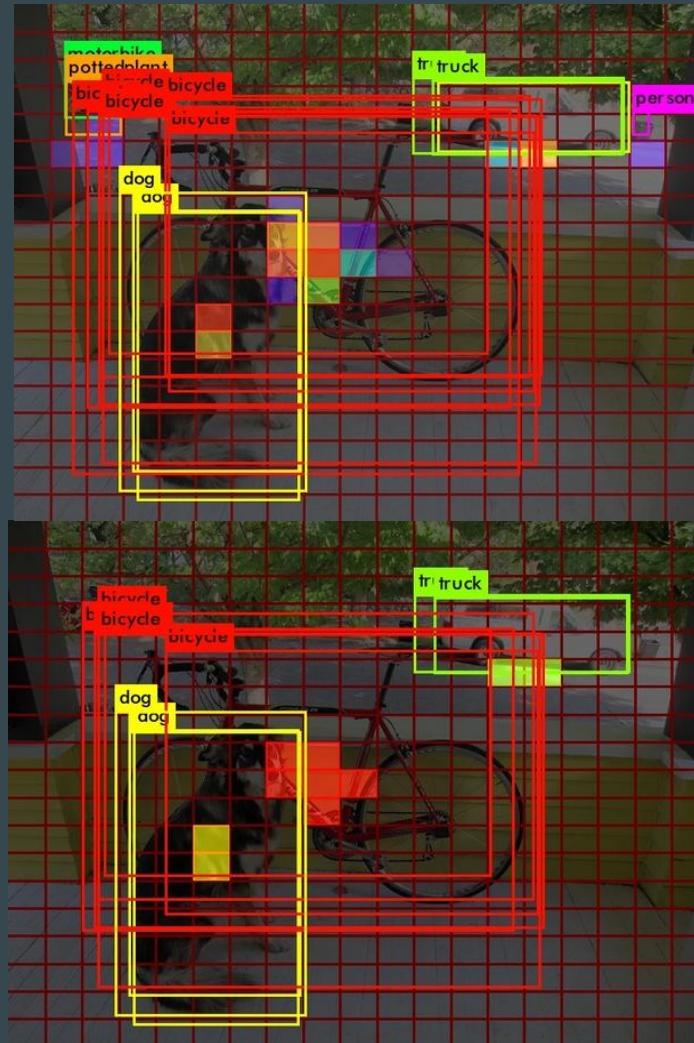
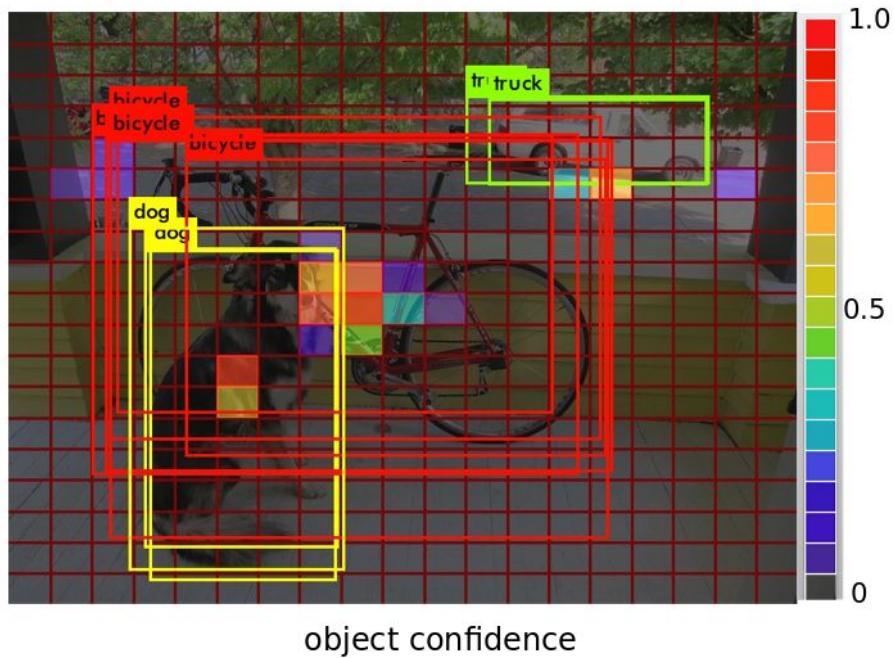
Yolo algorithm - beginning



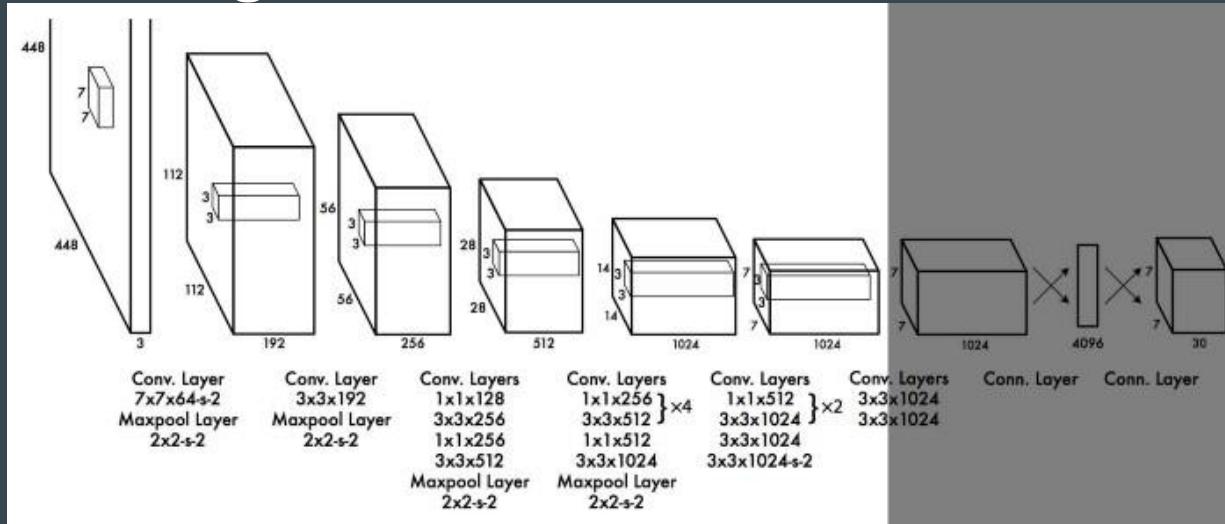
Yolo algorithm v2 - anchors



Grid visualization



Yolo algorithm - architecture - from 19 to 53



Type	Filters	Size	Output
Convolutional	32	3×3	256 x 256
Convolutional	64	$3 \times 3 / 2$	128 x 128
1x Residual	32	1×1	128 x 128
Convolutional	64	3×3	
Convolutional	128	$3 \times 3 / 2$	64 x 64
2x Residual	64	1×1	64 x 64
Convolutional	128	3×3	
Convolutional	256	$3 \times 3 / 2$	32 x 32
8x Residual	128	1×1	32 x 32
Convolutional	256	3×3	
Residual	512	$3 \times 3 / 2$	16 x 16
Convolutional	256	1×1	
8x Residual	512	3×3	16 x 16
Convolutional	1024	$3 \times 3 / 2$	8 x 8
Convolutional	512	1×1	
4x Residual	1024	3×3	8 x 8
Avgpool		Global	
Connected		1000	
Softmax			

Table 1. Darknet-53.

```
[convolutional]
batch_normalize=1
size=3
stride=1
pad=1
filters=256
activation=leaky
```

```
[convolutional]
size=1
stride=1
pad=1
filters=255
activation=linear
```

```
[yolo]
mask = 0,1,2
anchors = 10,13, 16,30,
33,23, 30,61, 62,45,
59,119, 116,90,
156,198, 373,326
classes=80
num=9
jitter=.3
ignore_thresh = .7
truth_thresh = 1
```

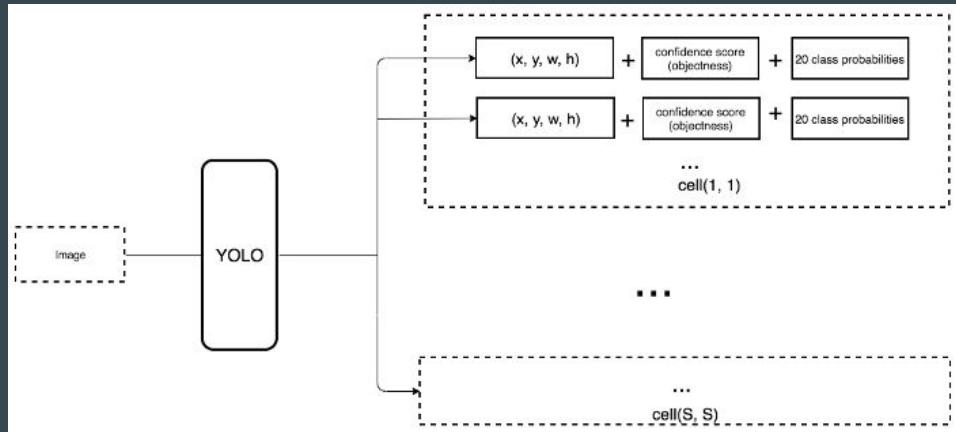
Bounding box parameters

The network has 3 output layers (13x13, 26x26, 52x52)

Each grid element has 3 bounding box proposals

The output layers of the network return 85 values:

- 4 bounding box coordinates
- 1 objectness (used for training)
- 80 confidences for the 80 classes



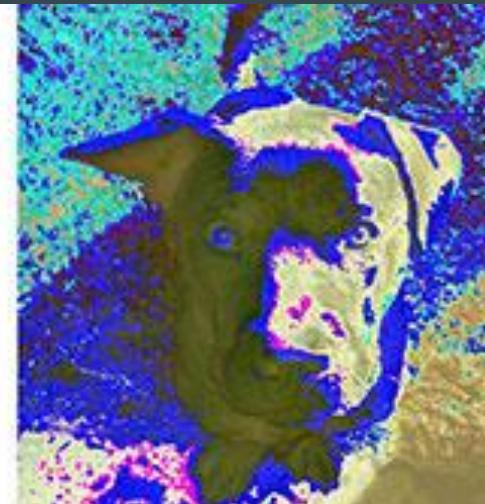
Preprocessing

- `blobFromImage`:
 - scale: scale images, pixels divide by 255 => in range 0...1
 - size
 - swapRB: OpenCV: BGR, normal image: RGB
 - mean: mean subtraction values => normalization

Blobs



$$\begin{aligned} R &= 124.96 \\ - G &= 115.97 = \\ B &= 106.13 \end{aligned}$$



FRAMEWORKS

- YOLOv3
- Python
- OpenCV
- COCO database: (common objects in context) from Microsoft. It is capable of detecting 80 common objects
- Test database: VOC 2012

Yolo implementation

- Setting parameters: confidence and threshold
- Preprocessing the images & frames with OpenCV
- Using Yolo handle images/ videos
- Implementing the evaluation class
- Evaluating results

Parameters

- **Confidence:** Minimum probability to filter weak detections
- **Threshold:** ignore bounding boxes that significantly overlap each other.

Confidence

Confidence = 0.5



Confidence

Confidence = 0.7



Threshold

Threshold = 0.1



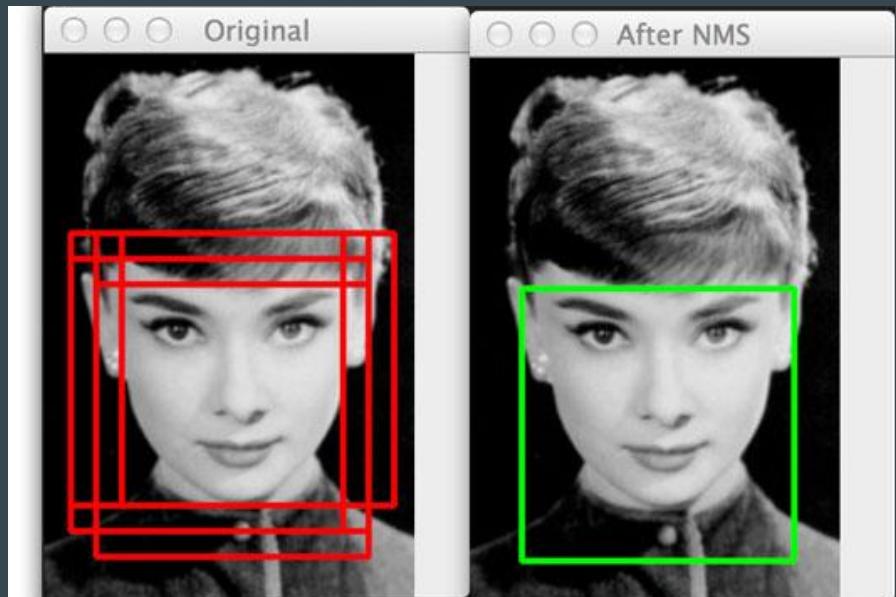
Threshold

Threshold = 0.7

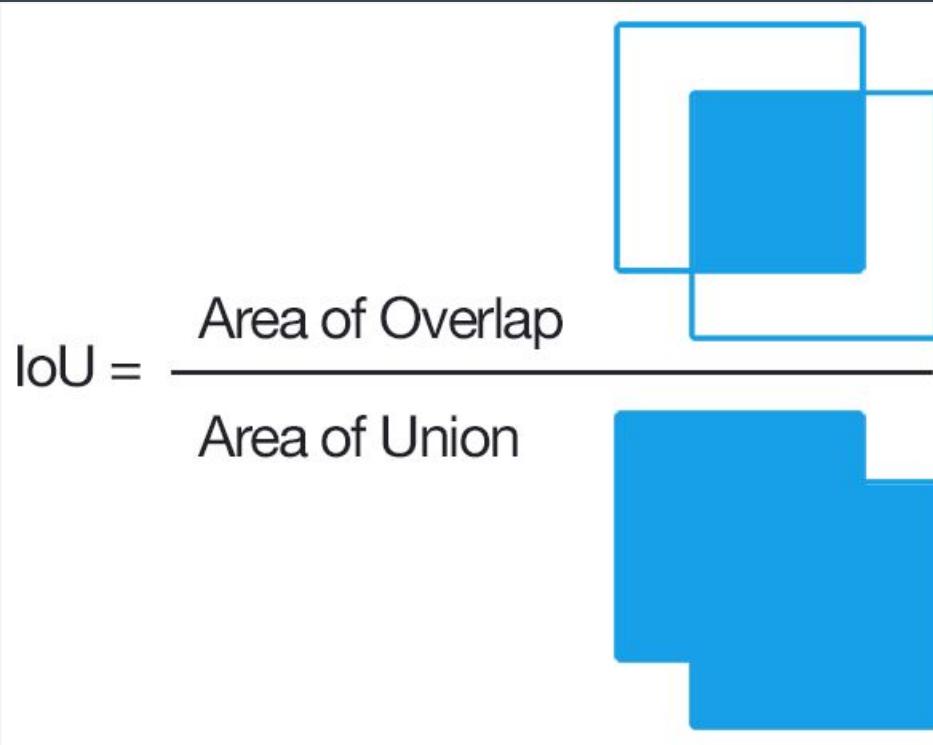


Non-Maximum Suppression

- Multiple bounding boxes
- Need to apply Non-Maximum Suppression
- Threshold high: multi-box
- Threshold low: one box



IoU: Intersection over Union



ACCURACY

- $Accuracy = \left(\frac{\sum i(IoU(x_i.t, x_i.p))}{|x|} + \left(\frac{\sum i(Label(x_i.t, x_i.p)))}{|x|} \right) / 3 + Diff(X.t, X.p) \right)$
- $Label(x, y) = \begin{cases} 1 & x = y \\ 0 & x \neq y \end{cases}$
- $Diff(x, y) = \left| \frac{|x| - |y|}{\max(|x|, |y|)} \right|$
- **X.t = actual classes**
- **X.p = prediction classes**

Test dataset building

Images and bounding boxes from VOC 2012:

General conditions for dataset:

- Image contains only labels that are also in the COCO dataset
- Image contains more than 5 objects

Datasets:

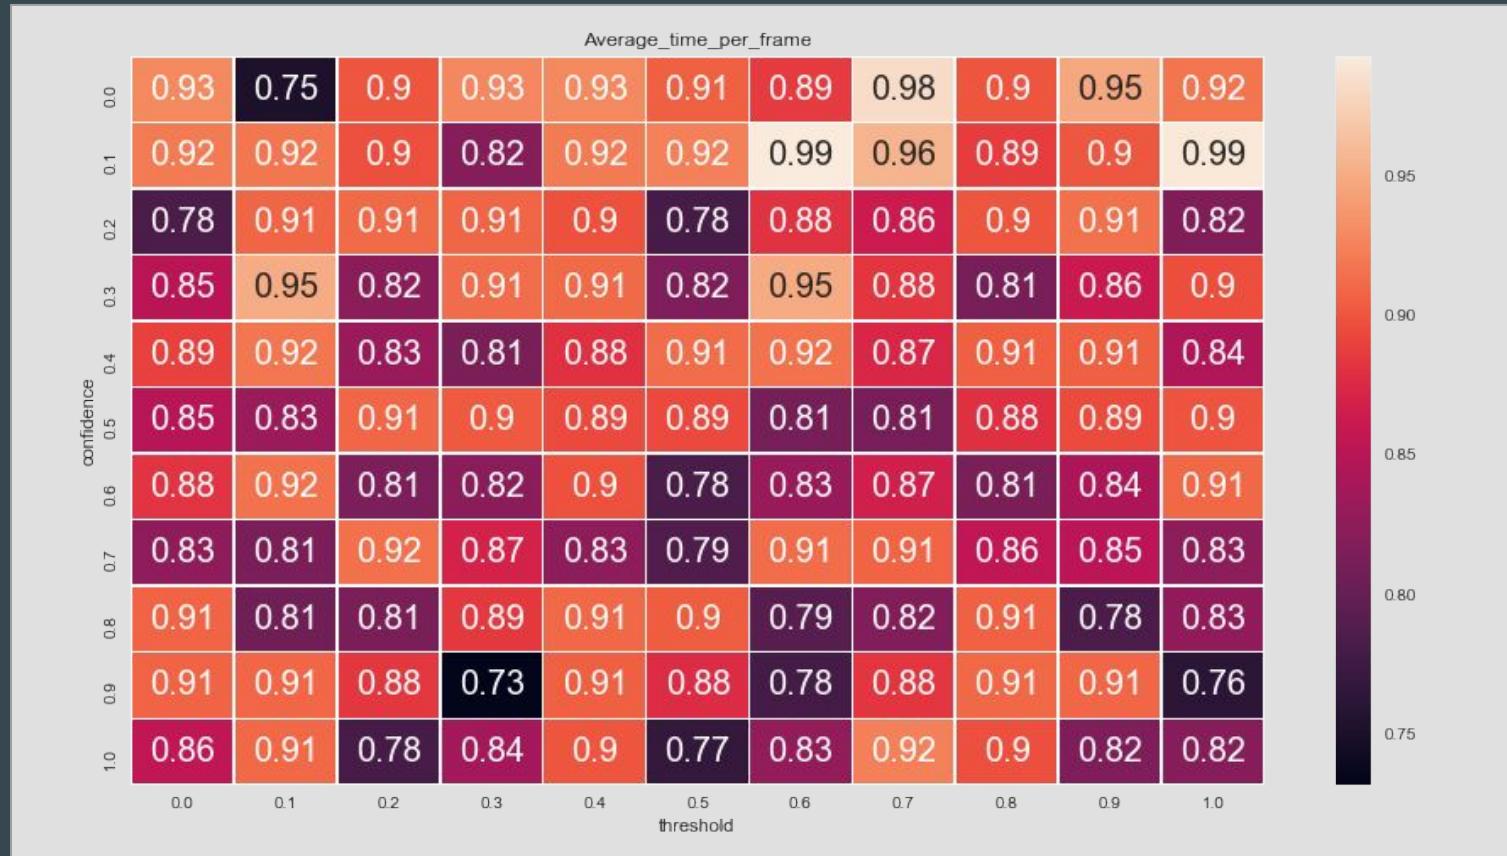
- Tiny - 20 images
- Small - 112 images
- Medium - 205 images
- Big - 1366 images

test/small/2011_001625.jpg
horse 154 96 420 279
horse 131 106 374 277
person 73 112 118 269
person 108 120 128 268
person 65 109 139 224
car 444 197 500 228
;

Accuracy graph



Average Time



Test results for best values

The best values seem to be:
confidence: 0.2 (or 0 or 0.1)
and threshold 0.6

Results:

- Tiny: 77.28%
- Small: 74.03%
- Medium: 73.87%
- Big: 73.44%

Results for base test
(conf: 0.5, thresh 0.3:

- Tiny: 65.33%
- Small: 66.63%
- Medium: 67.7%
- Big: 67.48%

Results for base test
(conf: 0.5, thresh 0.3:

- Tiny: 76.23%
- Small: 73.05%
- Medium: 72.88%
- Big: 72.38%

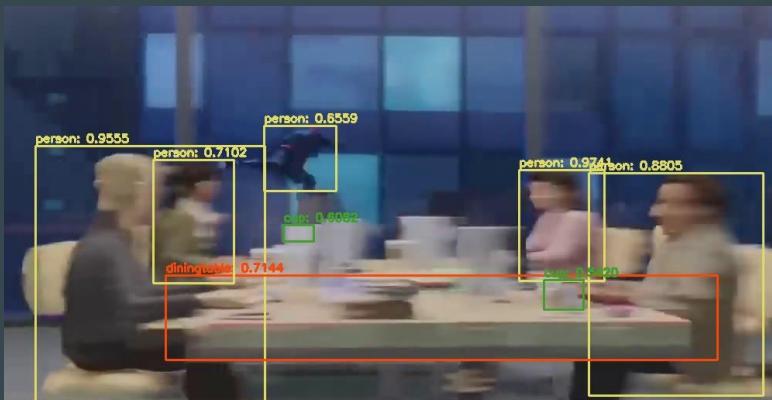
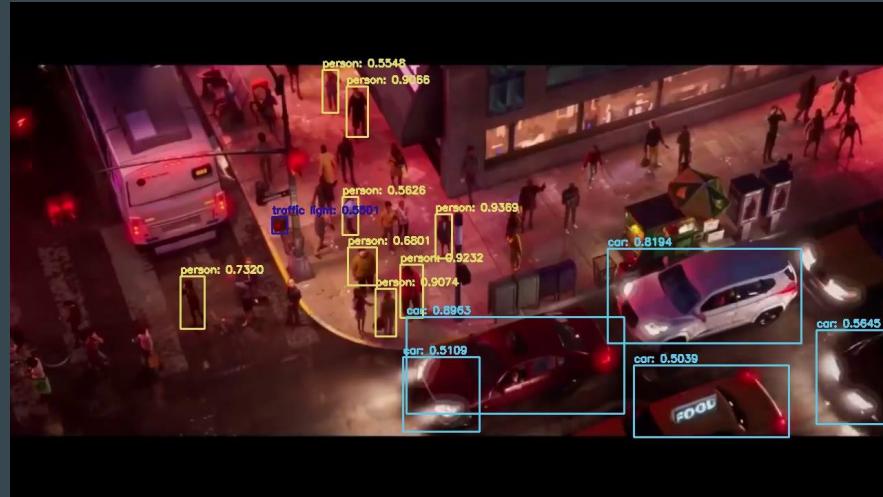
Speed (different laptops)

350 frames

- i7 4700MQ CPU @ 2.40GHz × 8
 - Total time: 180 s
 - Average time: 0.514 seconds per frame
- 2.7 GHz Intel Core i5
 - Total time: 290s
 - Average time: 0.829 seconds per frame
- i5 4200U CPU @ 1.6GHZ x 4
 - Total time: 339 s
 - Average time: 0.96 seconds per frame

Demo movies

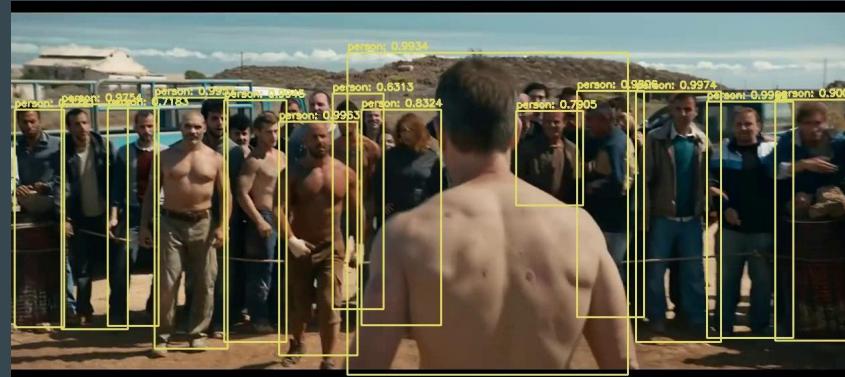
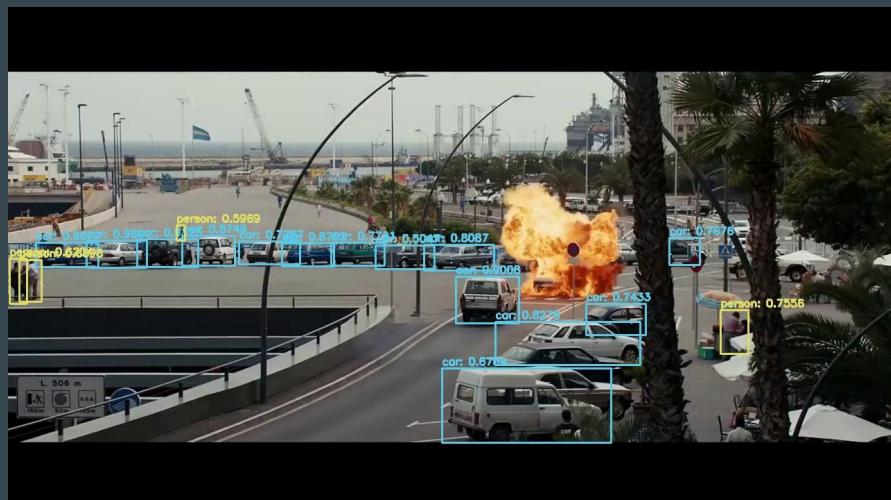
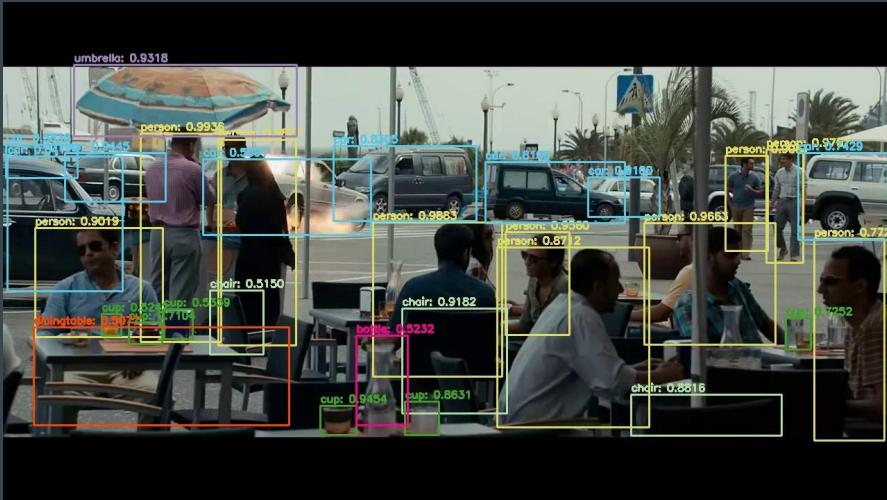
The good



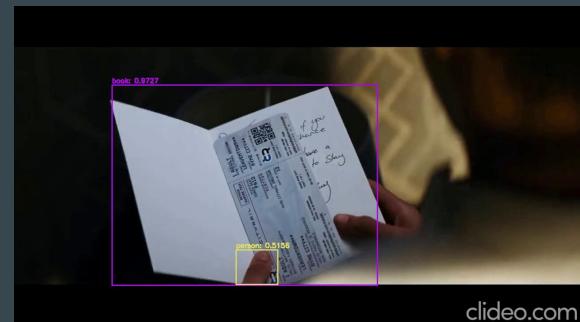
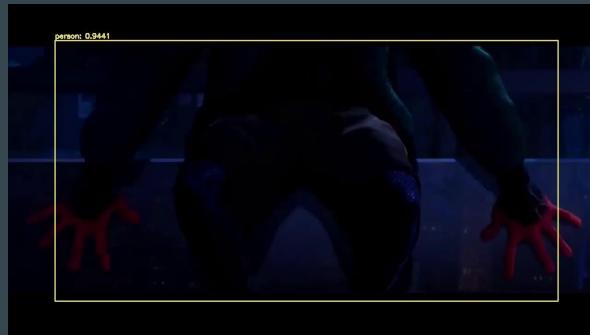
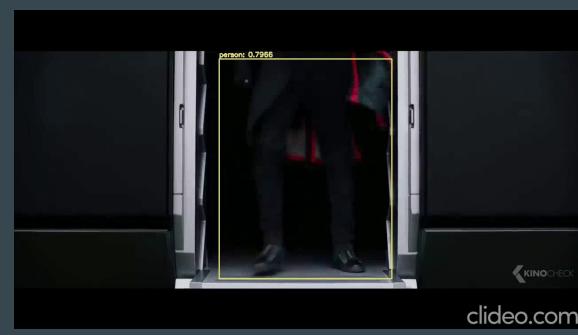
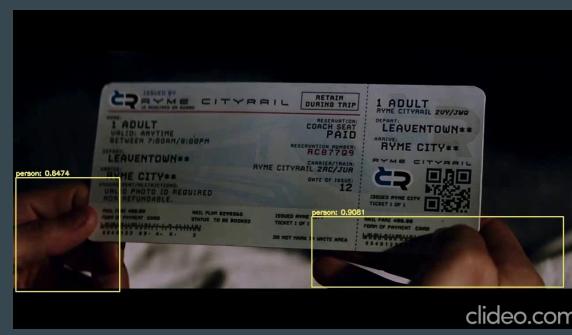
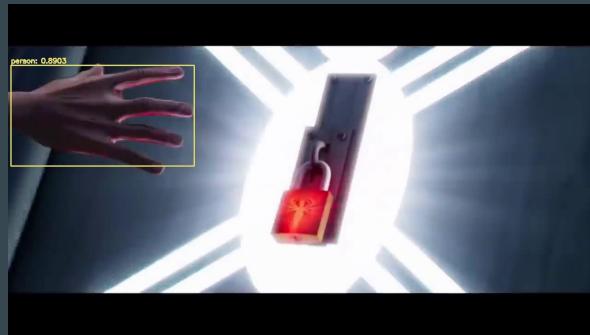
The bad



The many



Human detection

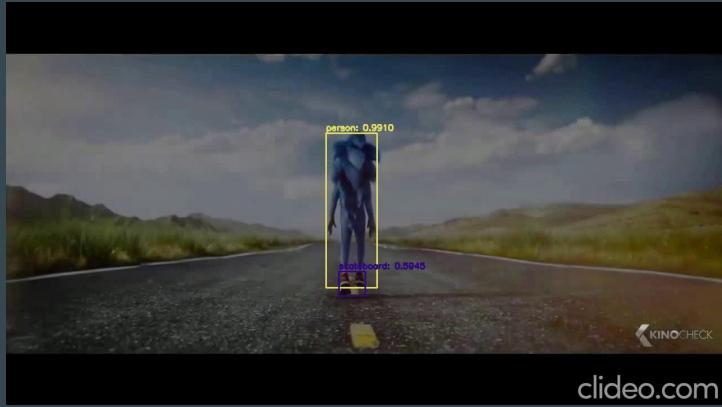


clideo.com

The fur of Pikachu



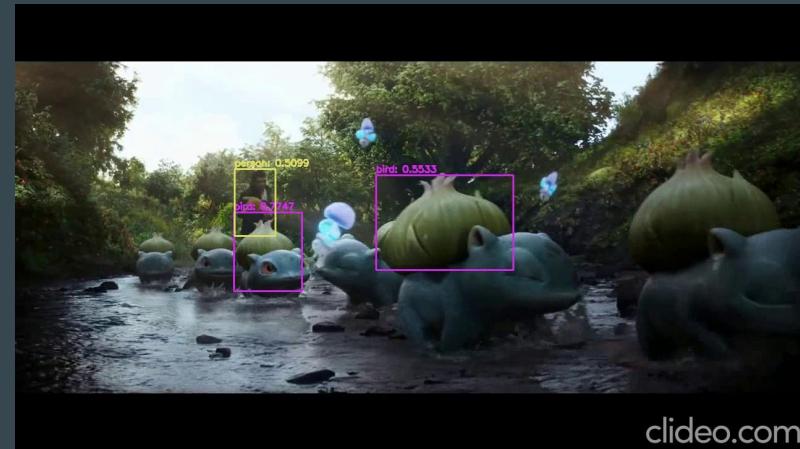
The pose of Sonic



Pokemon



clideo.com



clideo.com



clideo.com



clideo.com

The air and sea



The invisible skateboard



Bibliography

- <https://pjreddie.com/darknet/yolo/>
- <https://pjreddie.com/media/files/papers/yolo.pdf>
- <https://arxiv.org/pdf/1612.08242.pdf>
- <https://pjreddie.com/media/files/papers/YOLOv3.pdf>
- <https://www.pyimagesearch.com/2018/11/12/yolo-object-detection-with-opencv/>
- https://medium.com/@jonathan_hui/real-time-object-detection-with-yolo-yolov2-28b1b93e2088
- <https://github.com/xueeinsteins/darknet-vis>
- <https://towardsdatascience.com/review-yolov2-yolo9000-you-only-look-once-object-detection-7883d2b02a65>
- <https://www.pyimagesearch.com/2017/11/06/deep-learning-opencvs-blobfromimage-works/>
- <http://host.robots.ox.ac.uk/pascal/VOC/voc2012/#testdata>
- <https://www.youtube.com> - (trailers)